

Something about machine learning

Episode 2


**Viviana Acquaviva
(CUNY)**

vacquaviva@citytech.cuny.edu

Let us focus on supervised learning for now.

Important:
**you shouldn't use all of your learning set
to build your model.**

**It is customary to split the learning set into
a **training set** and **test set****



**By building a model on the training set and applying it
to the test set, you “mimic” what happens when your model sees new data
for which the labels are not known.**

Otherwise you would be too optimistic!

Note:
You still are.

For example...

Math Quiz #1 - Teacher's Answer Key

1) $2 \ 4 \ 5 = 3$

2) $5 \ 2 \ 8 = 2$

3) $2 \ 2 \ 1 =$

4) $4 \ 2 \ 2 =$

TEST

**Diagnosing and Improving
Machine Learning
algorithm performance:
cross validation,
performance metrics,
diagnostics**

The goal of the training set and test set split is to be able to evaluate performance on unseen examples.

The test set “mimics” new data.

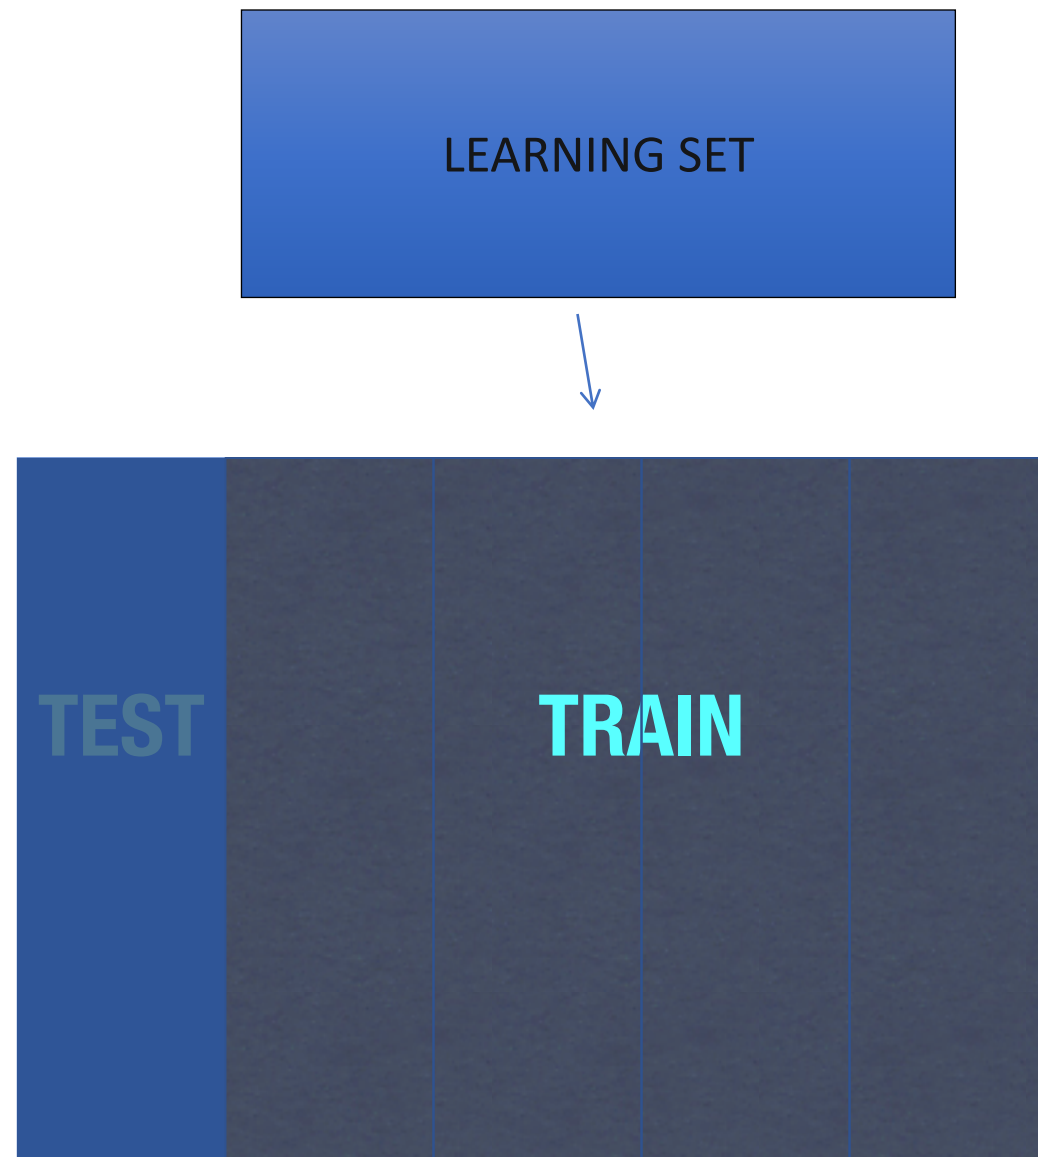
However, it might be better to pick more than one test split.

Why would this be a good idea?

1. We use all the training data for training!

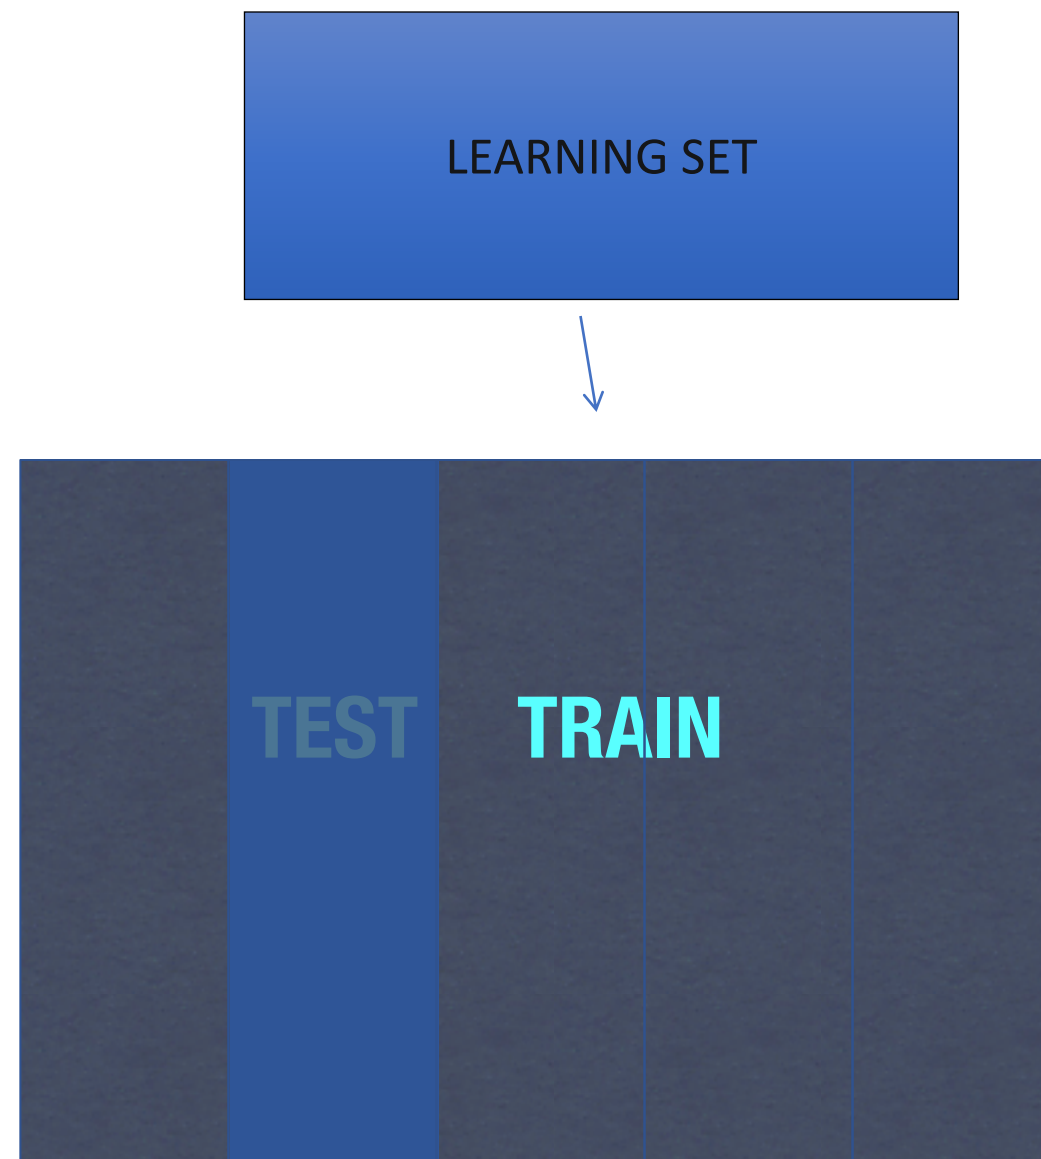
2. We avoid the risk of under/overestimating performance because of a “weird” pick of train/test split.

k-FOLD CROSS VALIDATION



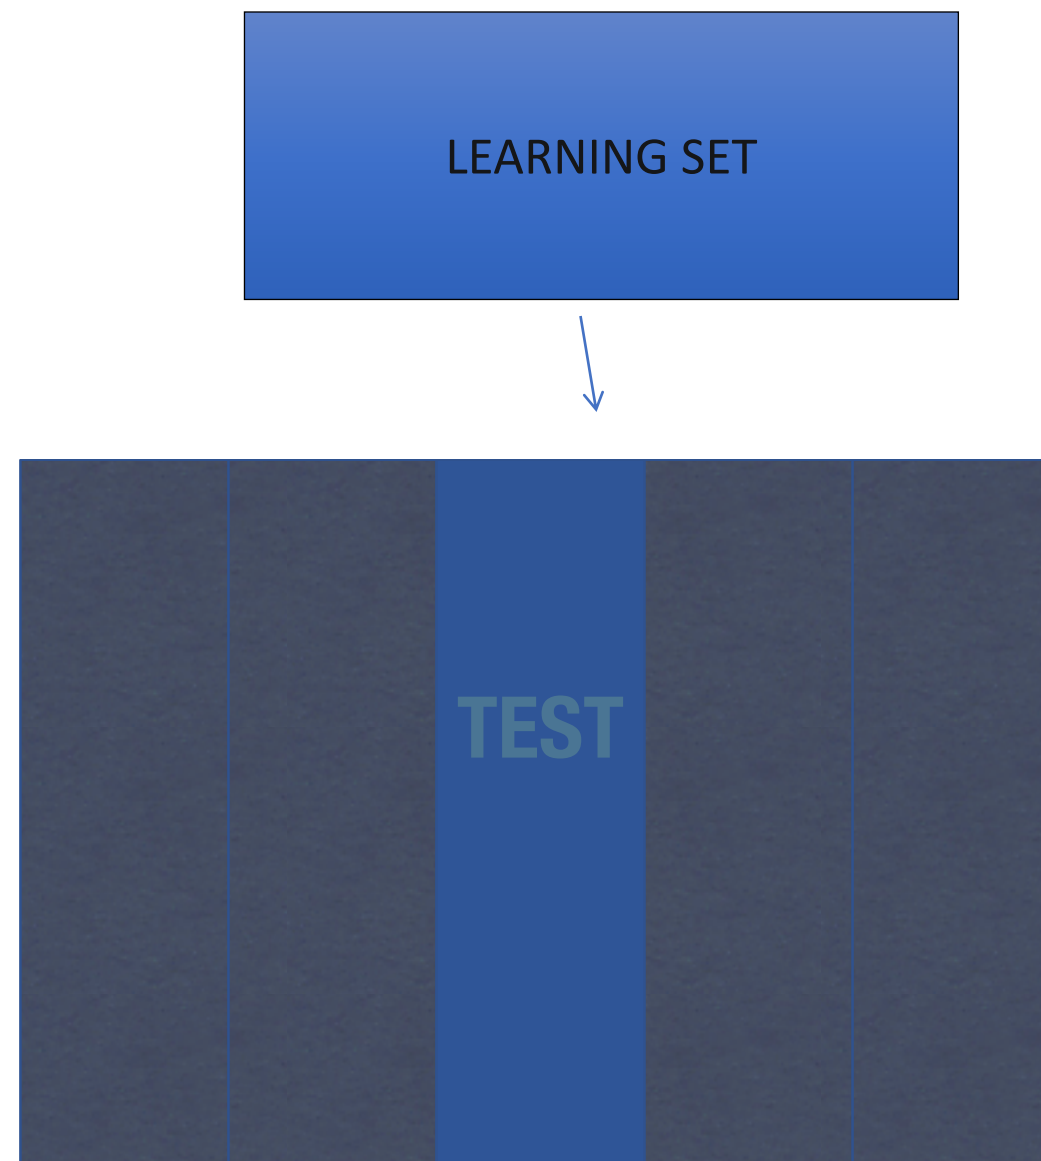
5 folds ($k = 5$)

k-FOLD CROSS VALIDATION



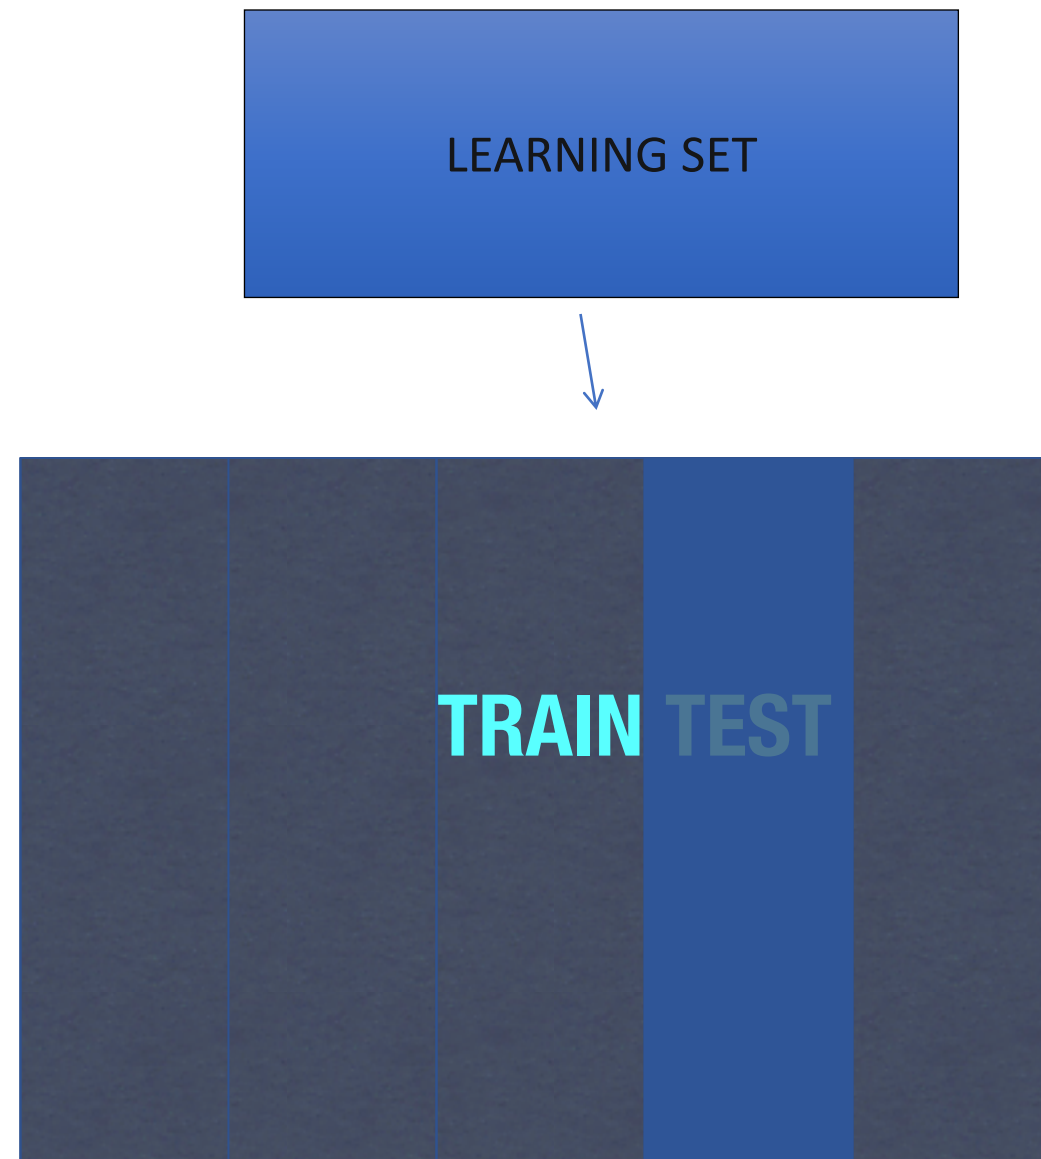
5 folds ($k = 5$)

k-FOLD CROSS VALIDATION



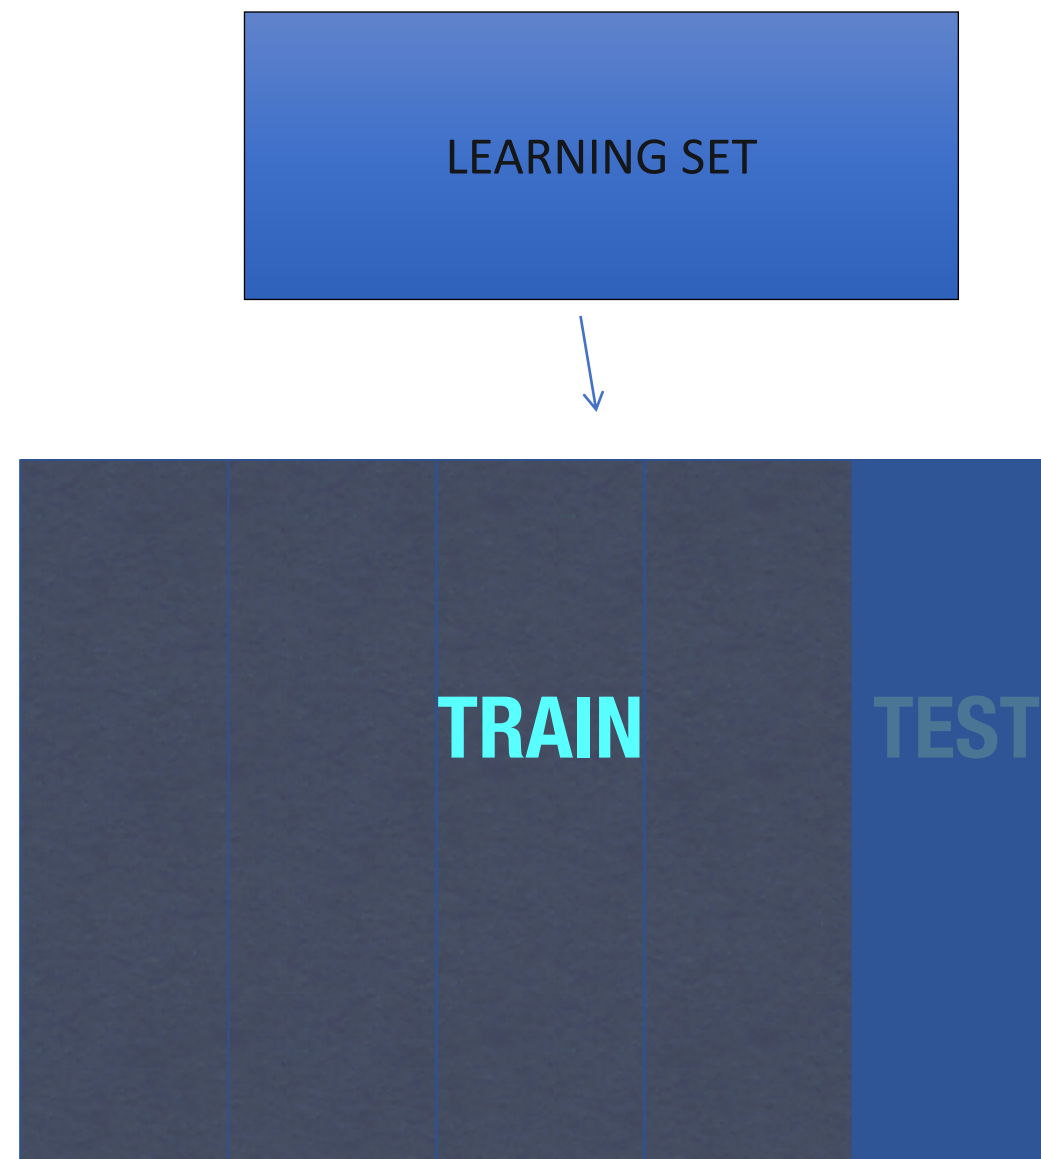
5 folds ($k = 5$)

k-FOLD CROSS VALIDATION



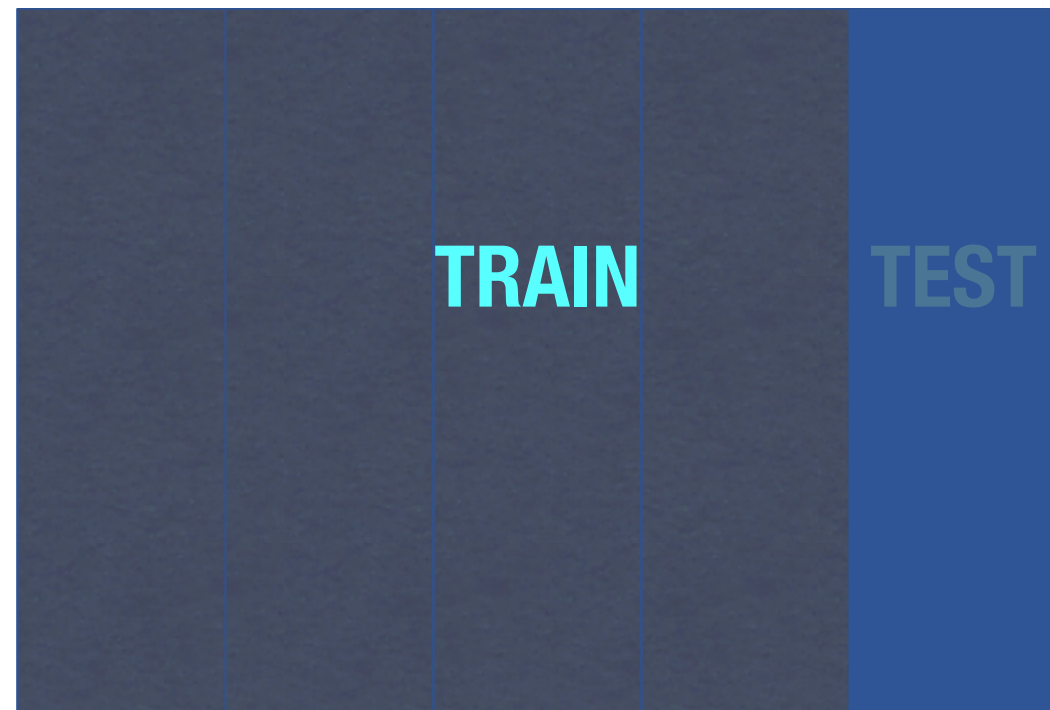
5 folds ($k = 5$)

k-FOLD CROSS VALIDATION



5 folds ($k = 5$)

k-FOLD CROSS VALIDATION



5 folds ($k = 5$)

- 😊 **We use all learning data**
- 😊 **The standard deviation of scores vector gives us an idea of average performance + uncertainty**
- 😞 **It takes 5x more time to run**

Diagnosing a ML algorithm

BIAS

Algorithm
can't capture
complexity
of rule
connecting
input and output


UNDERFITTING

VARIANCE

Algorithm
is excessively
tailored
to training set
and generalizes
poorly

OVERFITTING


Bias/variance



$$\theta_0 + \theta_1 x$$

High bias
(underfit)


$$d=1$$



$$\theta_0 + \theta_1 x + \theta_2 x^2$$

“Just right”

$$d=2$$





$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

High variance
(overfit)

$$d=4$$

How can we diagnose high variance vs high bias?

figure from Jake VanderPlas' book



High bias: test and train error are similar but high

High variance: there is a gap between test and train error because algorithm does not generalize well

Improving high bias

- 1. Try using different features.**
- 2. Try engineering new features.**
- 3. Try a more complex algorithm.**


Improving high variance

- 1. Try reducing the number of features.**
- 2. Try a less complex algorithm/change parameters.**

Also: Check if you need more training data

Useful diagnostics: learning curves

plot performance of algorithm for train and test set
as a function of size of training set



Note: which algorithm is the best?

- High bias low variance**
- High variance low bias**
- Lowest gap between train/test**
- Highest test scores**

Note: which algorithm is the best?

- High bias low variance**
- High variance low bias**
- Lowest gap between train/test**
- Highest test scores**

DECISION TREES

- Work by splitting data on different values of features
- Simplest trees are binary trees
- If categorical features, the split would be on yes/no
- If numerical, the split would be on a certain value (e.g. $x > 100$ or $x < 100$)



Example: Look at this 2-feature data set. How should we split?




Figure credit:
Gilles Louppes




Figure credit:
Gilles Louppes

Should we stop?






Figure credit:
Gilles Louppes

Decision trees: defined by splits and leaves



Figures credit:
Gilles Louppes

How many splits in this tree?

How many leaves?

How do we decide whether we should keep splitting?

Pseudo code for decision trees

```

function BuildDecisionTree(L)
  Create node  $t$  from the learning sample  $L_t = L$ ;
  calculate (im)purity
  if the stopping criterion is met for  $t$  then
     $y^{\wedge} =$  some constant value/class (MAKE PREDICTION)
  else
    Find the split on  $L_t$  that maximizes impurity
    decrease
     $s^* = \arg \max_{s \in Q} \Delta i(s, t)$ 
    Partition  $L_t$  into  $L_{t_L} \cup L_{t_R}$  according to  $s^*$ 
     $t_L = \text{BuildDecisionTree}(L_L)$ 
     $t_R = \text{BuildDecisionTree}(L_R)$ 
  end if
  return  $t$ 
end function

```

Code adapted from Gilles Louppes

stopping criterion
Gini (im)purity = 0

Gini (node L) =

$$1 - \sum f(i)^2$$

where $f(i)$ is the frequency of
the i -th class

Gini (splits L_t and L_r) =

$$\frac{L_L}{L} * (1 - \sum f(i)^2) + \frac{L_R}{L} * (1 - \sum f(i)^2)$$

where $f(i)$ is the frequency of
the i -th class

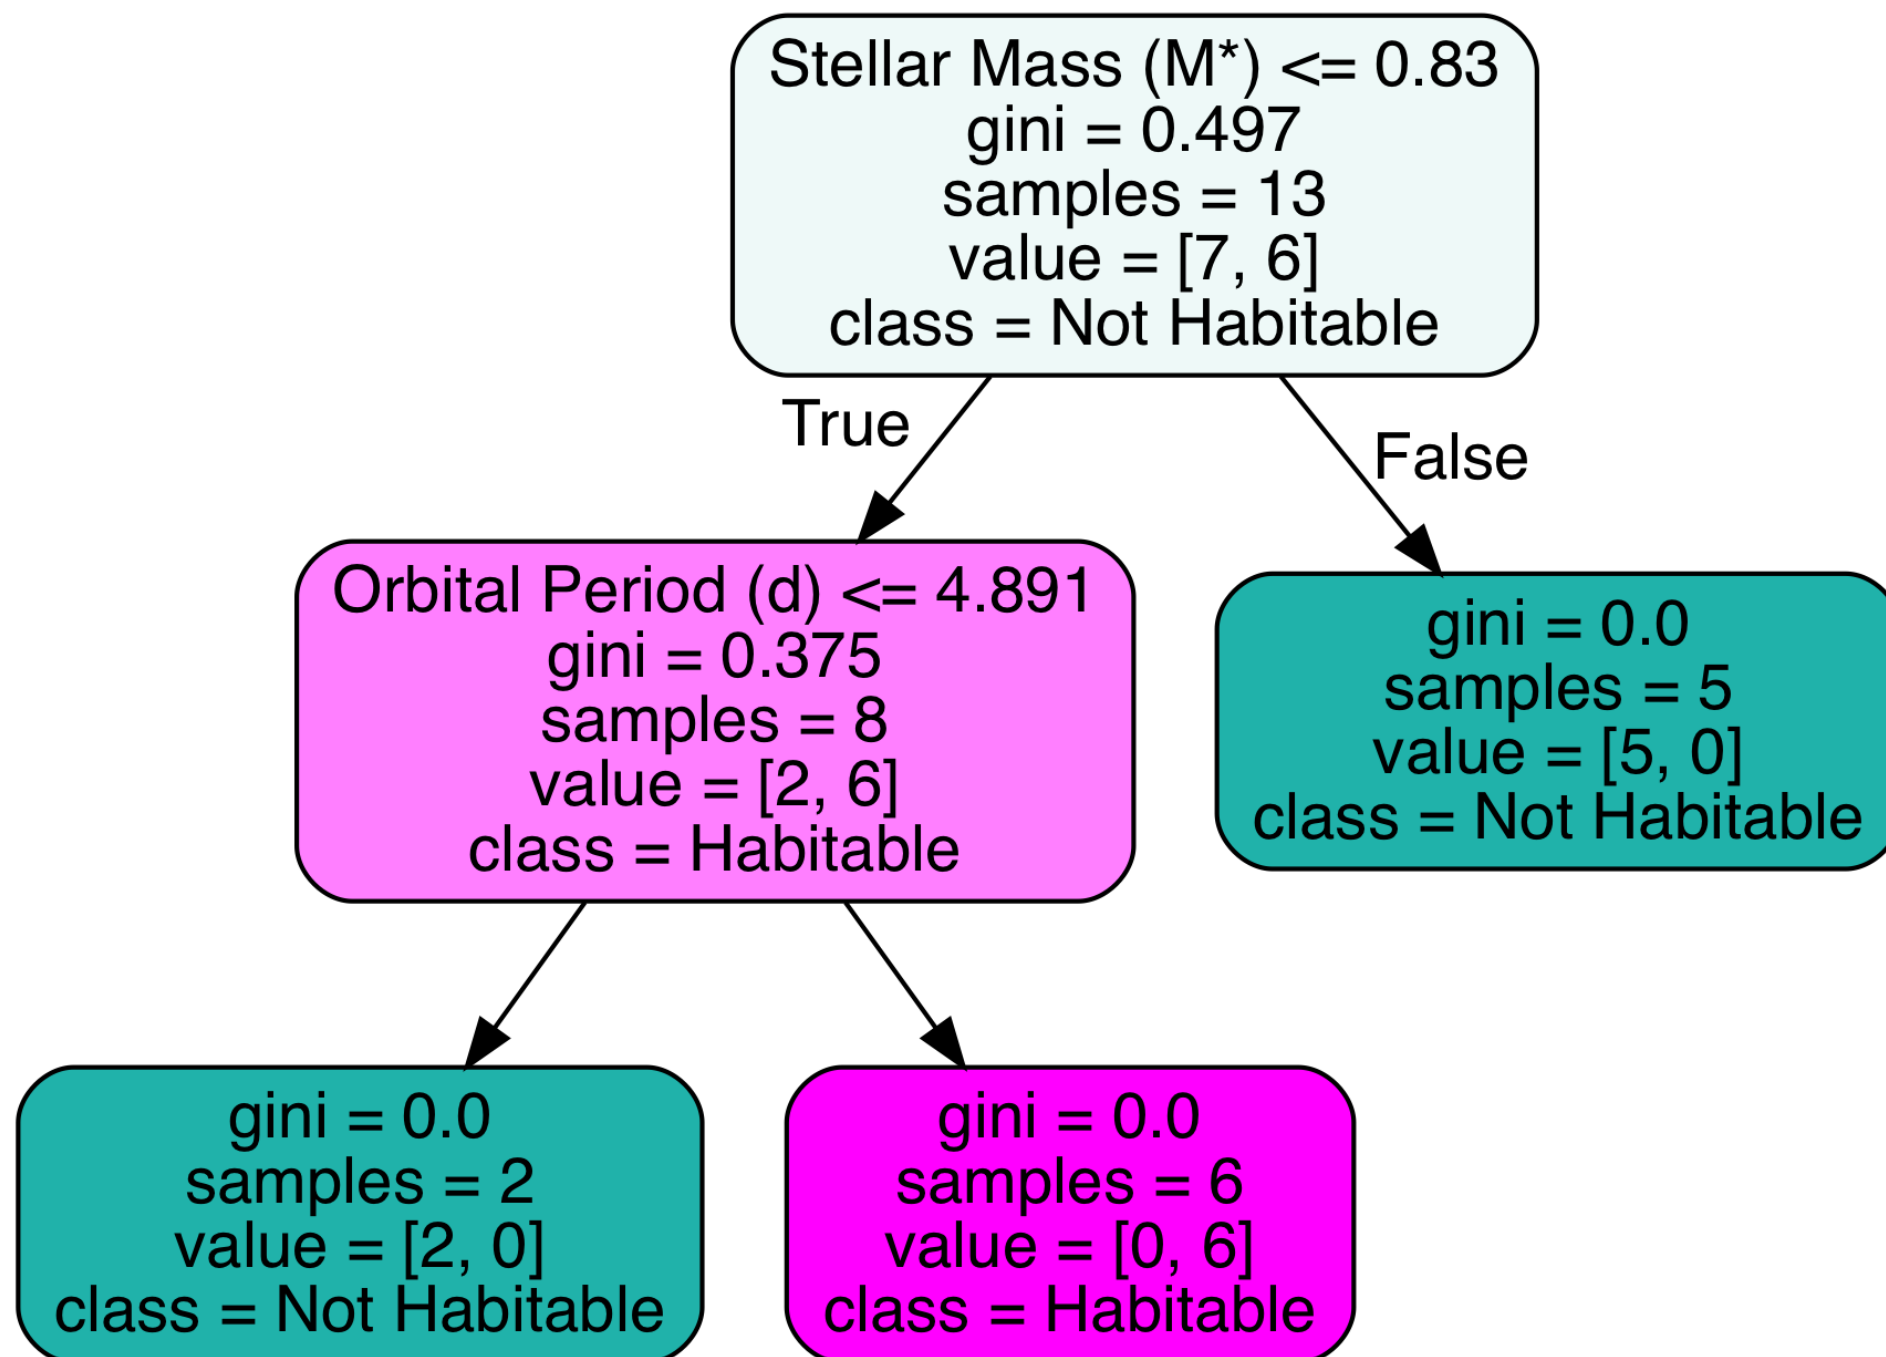
Note: **splits**
happen along features!

Our first example will be a supervised classification problem, in which we are trying to decide if a planet is habitable, based on its distance from parent star, mass of parent star, and orbital period.

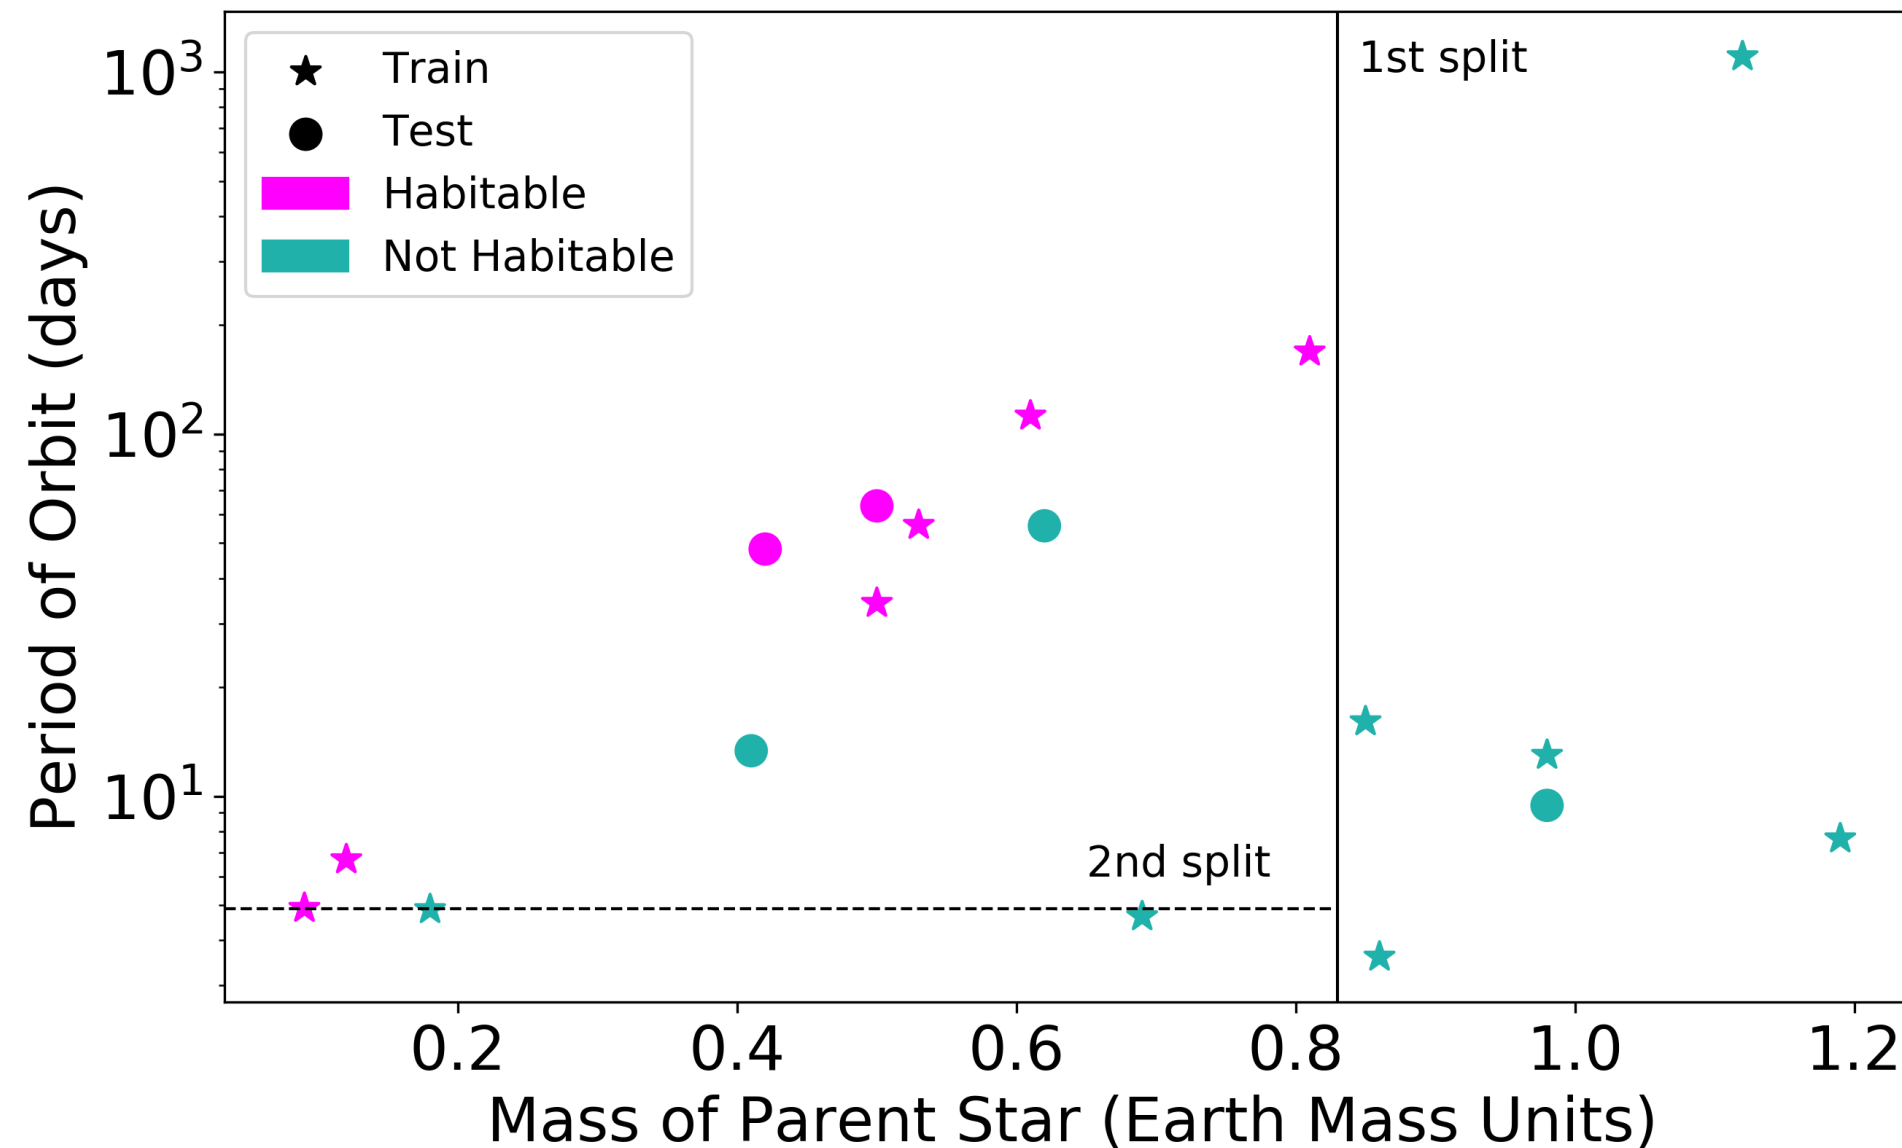
NAME	Stellar Mass (M_{\odot})	Orbital Period (days)	Distance (AU)	Habitable?
Kepler-736 b	0.86	3.60	0.0437	0
Kepler-636 b	0.85	16.08	0.1180	0
Kepler-887 c	1.19	7.64	0.0804	0
Kepler-442 b	0.61	112.30	0.4093	1
Kepler-772 b	0.98	12.99	0.1074	0
Teegarden's Star b	0.09	4.91	0.0252	1
K2-116 b	0.69	4.66	0.0481	0
GJ 1061 c	0.12	6.69	0.035	1
HD 68402 b	1.12	1103	2.1810	0
Kepler-1544 b	0.81	168.81	0.5571	1
Kepler-296 e	0.5	34.14	0.1782	1
Kepler-705 b	0.53	56.06	0.2319	1
Kepler-445 c	0.18	4.87	0.0317	0
HD 104067 b	0.62	55.81	0.26	
GJ 4276 b	0.41	13.35	0.0876	
Kepler-296 f	0.5	63.34	0.2689	
Kepler-63 b	0.98	9.43	0.0881	
GJ 3293 d	0.42	48.13	0.1953	

Table 2.1: Learning set for the habitable planets problem.

Let's see what sklearn says

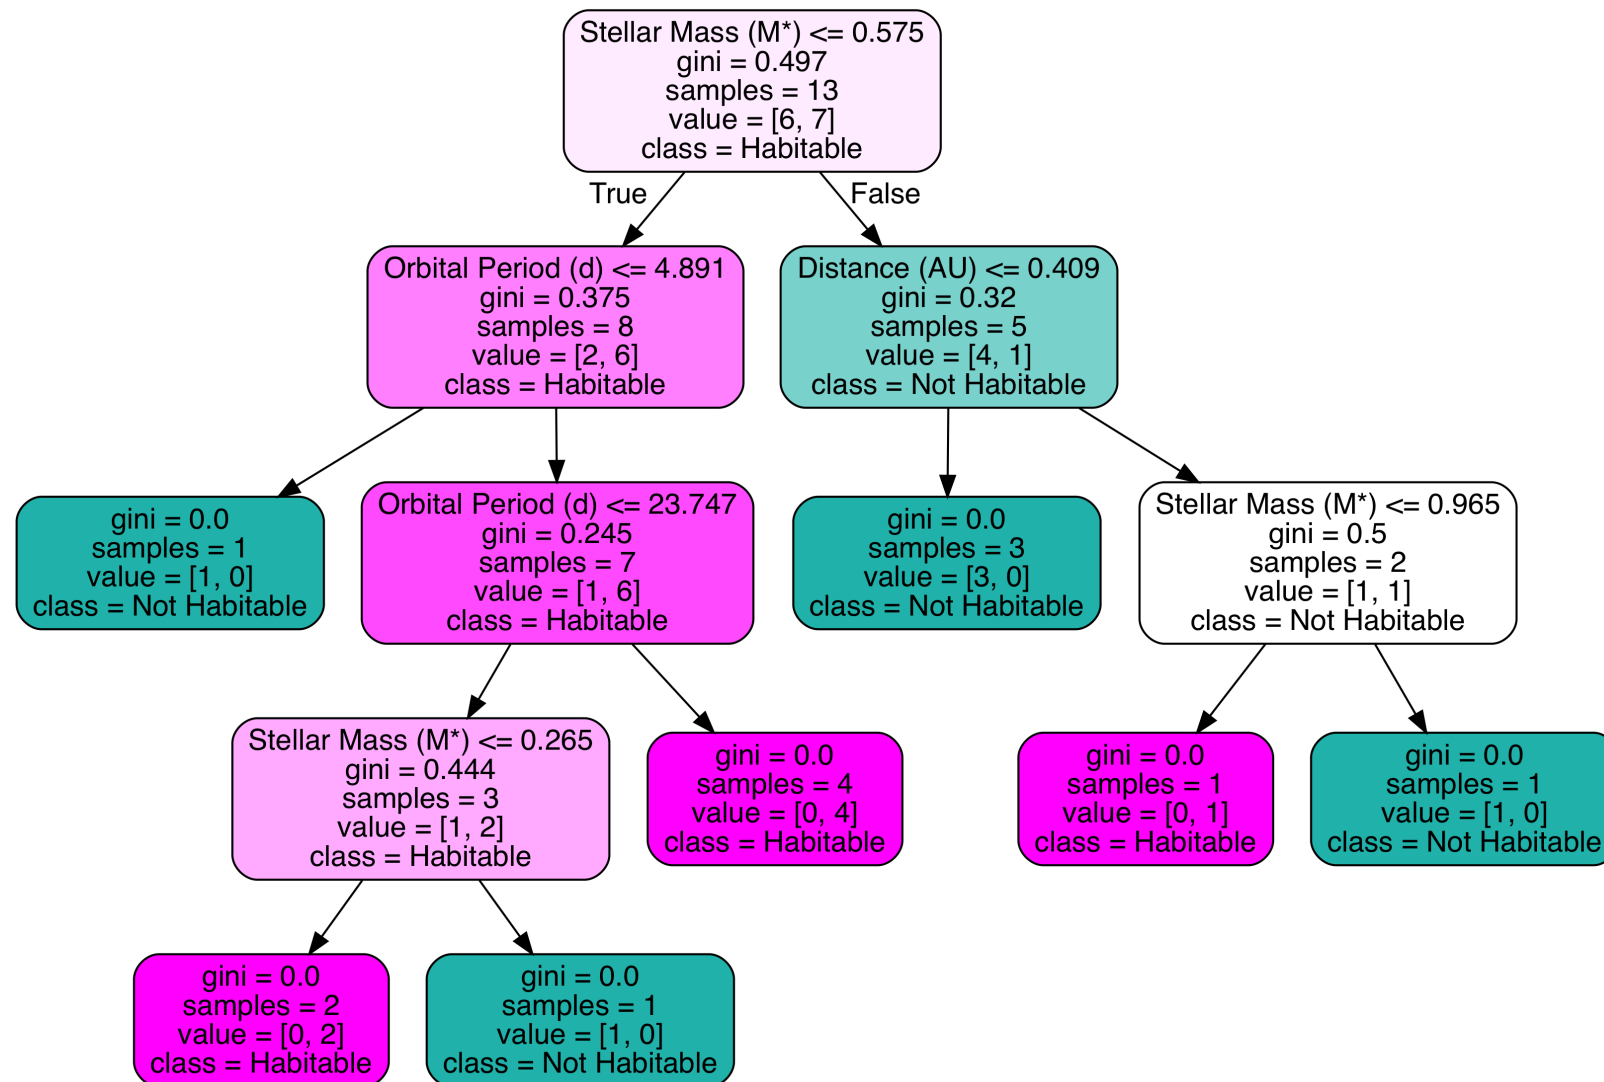


And visualize the criteria that we found!




Can you figure out the accuracy on the test set?

**Note (but you have to take my word for it):
if you use the last 13 rows for training and the first
5 for testing, you get this tree:**

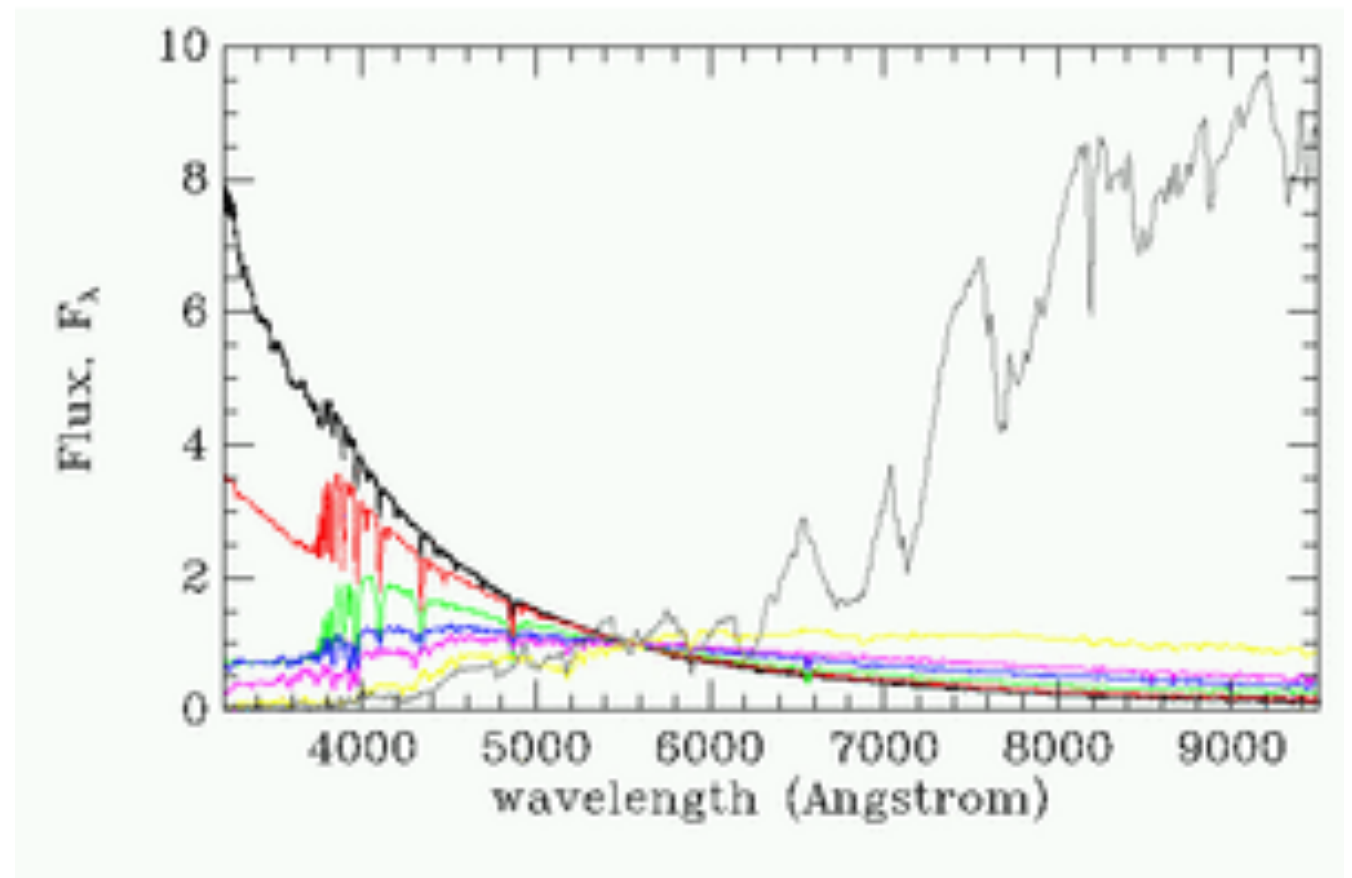


and 100% accuracy on test set (remember our instability problem?)

Coding exercise: supervised classification problem, in which we are trying to decide if a star is a variable star (RR Lyrae), based on imaging data in 5 bands (u, g, r, i, z) and four colors (u-g, g-r, r-i, i-z)



range of observed brightness



spectra of different stellar types